

Module 1 Reflection | Dan Johnson | 6/22/26

What assumption about AI did this module confirm, challenge, or complicate? Describe the assumption and what changed.

Before doing this reading, I was unaware of the existence of foundation models and their effects on generative AI as we know and use it today. I had assumed that all the major LLMs had been developed independently rather than being built on a common foundation, which is apparently incorrect. I think that helps me to better understand the concerns that people have about bias in LLMs: if they're all coming from the same source, then any issue of bias in the foundation will also be replicated in models relying on that foundation (homogenization).

What one aspect of how LLMs work technically surprised you most — and why does it matter for your professional context?

The aspect that surprised me the most was the concept of emergence: that AI models can be used for purposes they weren't trained for and actually be reasonably good at those tasks. That means, though, that there might be a ceiling on their performance on any given task. I think that impacts my work in that depending on what kind of task we're trying to get AI to help with, we may be able to accomplish it with an off-the-shelf LLM, we might have to look into getting access to a more specialized tool, or we might have to work on developing a tool ourselves.

Name one specific recurring task in your role that you now think could be affected by current AI. Be concrete: name the task, identify the AI mechanism, and describe one realistic impact scenario — positive or cautionary.

The first task that comes to mind when considering AI integration is processing student feedback from our end-of-semester and individual session student feedback. I think it's important for faculty to read through themselves, but it would also help to have an at-a-glance version that would be easier to review and provide additional perspective since it can be easy to focus only on constructive feedback. This would likely be possible through existing LLM resources - the trick would be figuring out what kind of format we want and then fine-tuning a prompt that can perform the task. If it were to work out well, we could

make it much easier for our course teams to make decisions on what changes to prioritize for next year's run of a course.