

Prompt	Gemini Output summary	Gemini accuracy	Gemini relevance	Gemini format	ChatGPT Output summary	ChatGPT accuracy	ChatGPT relevance	ChatGPT format
What year was the Transformer architecture first introduced in a peer-reviewed paper, and what was that paper's exact title? Please provide an APA citation.	The Transformer architecture was first introduced in 2017 (Sajun et al., 2024; Vaswani et al., 2017). The exact title of the peer-reviewed paper is: "Attention Is All You Need" (Vaswani et al., 2017).	4	5	5	The Transformer architecture was first introduced in 2017 in the peer-reviewed conference paper Attention Is All You Need by Ashish Vaswani and seven coauthors. The paper proposed the Transformer, a neural network architecture based entirely on attention mechanisms, replacing the recurrent and convolutional structures that had previously dominated sequence modeling. The paper was presented and published in the proceedings of the 31st annual conference on Advances in Neural Information Processing Systems 30 (NeurIPS 2017) in December 2017.	5	4	3

Summarize the following abstract in 100 words for a non-specialist audience: Making language models bigger does not inherently make them better at following a user's intent. For example, large language models can generate outputs that are untruthful, toxic, or simply not helpful to the user. In other words, these models are not aligned with their users. In this paper, we show an avenue for aligning language models with user intent on a wide range of tasks by fine-tuning with human feedback. Starting with a set of labeler-written prompts and prompts submitted through a language model API, we collect a dataset of	Simply making AI language models larger doesn't automatically make them better at following directions; they can still generate unhelpful, toxic, or untruthful answers. To fix this, researchers developed InstructGPT by training OpenAI's GPT-3 using human feedback. First, humans showed the AI how to properly respond to prompts. Then, humans ranked different AI responses, teaching the model which answers were best through reinforcement learning. The results were dramatic: a smaller InstructGPT model was preferred over a standard GPT-3 model that was 100 times larger. It also proved to be more	5	5	5	This paper argues that simply making language models larger does not ensure they better understand or follow what users want. To improve alignment, the researchers trained GPT-3 using human feedback. First, people provided examples of good responses, and then they ranked different model outputs. These rankings were used to further train the model through reinforcement learning. The resulting system, called InstructGPT, produced responses that people preferred over those from the much larger GPT-3 model, despite being 100 times smaller. InstructGPT was also more truthful and less likely to generate toxic content, showing	5	5	5
--	--	---	---	---	---	---	---	---

labeler demonstrations of the desired model behavior, which we use to fine-tune GPT-3 using supervised learning. We then collect a dataset of rankings of model outputs, which we use to further fine-tune this supervised model using reinforcement learning from human feedback. We call the resulting models InstructGPT. In human evaluations on our prompt distribution, outputs from the 1.3B parameter InstructGPT model are preferred to outputs from the 175B GPT-3, despite having 100x fewer parameters. Moreover, InstructGPT models show improvements in

truthful and less toxic, proving human feedback is key to aligning AI with human intent.

that human feedback can significantly improve AI behavior.

truthfulness and reductions in toxic output generation while having minimal performance regressions on public NLP datasets. Even though InstructGPT still makes simple mistakes, our results show that fine-tuning with human feedback is a promising direction for aligning language models with human intent.								
Write a Python function that accepts a list of numbers and returns a dictionary with keys mean, median, and standard_deviation. Include a docstring and one usage example.	'mean': 7, 'median': 7, 'standard_deviation': 4.3012	5 (used sample standard deviation)	5	5	'mean': 7.0, 'median': 7, 'standard_deviation': 3.847076812334269	5 (used population standard deviation)	5	5
Draft a professional email declining a speaking invitation due to a scheduling conflict. The tone	Thank you so much for inviting me to speak at the upcoming [Event Name] on [Date]. I	5	5	3	Thank you very much for inviting me to speak at [Event Name]. I am honored by the	5	5	5

should be warm and leave the door open for future invitations. Length: 150–200 words.	am incredibly honored that you reached out, and the agenda looks fantastic. Regrettably, I must decline the invitation due to a prior commitment that creates a scheduling conflict on that day. Because of this, I won't be able to give your event and audience the full attention they deserve. I am genuinely disappointed to miss out this time around, as I deeply admire the work you are doing with [Organization/Event Name]. I would love to stay in touch and hope you will keep me in mind for future panels, conferences, or speaking opportunities down the road. Wishing you and your team a highly				opportunity and appreciate your consideration. Unfortunately, due to a prior commitment that conflicts with the event schedule, I will be unable to participate this time. I am disappointed to miss the chance to contribute and engage with your audience, as the event sounds both meaningful and well organized. I am grateful for the invitation and for the work you and your team are doing. I hope the event is a great success and that the attendees find it valuable and inspiring. Please do keep me in mind for future speaking opportunities. I would welcome the chance to collaborate when schedules align, and I would be happy to			
--	--	--	--	--	---	--	--	--

	successful and inspiring event! Warm regards,				explore ways to participate in upcoming events. Thank you again for thinking of me. I wish you all the best with the event and look forward to staying in touch. Warm regards,			
Write a two-sided argument — two sentences pro, two sentences con — on the following claim: Universities should require disclosure of all AI use in student assignments.	Pro-Disclosure Requiring students to disclose all AI use ensures academic integrity by providing transparency into how a piece of work was actually produced. This clear baseline allows educators to fairly evaluate a student's independent critical thinking and research skills rather than grading the capabilities of an algorithm. Anti-Disclosure Enforcing a blanket disclosure policy is practically impossible due to the deep integration	4.5	5	5	Pro: Requiring disclosure of all AI use promotes academic transparency and helps instructors evaluate how students developed their work. It can also encourage responsible AI use by making students accountable for when and how they rely on these tools. Con: A universal disclosure requirement may be difficult to enforce and could create uncertainty about what counts as AI assistance. It may also discourage	5	5	5

	of AI features within everyday tools like spell-checkers, word processors, and search engines. Furthermore, it risks penalizing students unfairly based on unreliable detection tools while creating an atmosphere of surveillance that stifles the natural adoption of modern tech.				beneficial uses of AI, such as brainstorming or proofreading, even when those uses do not undermine learning.			
--	--	--	--	--	---	--	--	--

Neither model appeared as a clear frontrunner. On three tasks, I rated them equally, Gemini performed better on one, and ChatGPT performed better on one. I was surprised that they interpreted standard deviation differently (though I know from my statistics course that both of these are valid measures, so it's not unreasonable). I think this speaks to the importance of being very clear when working with an LLM on **exactly** what you want from it.