

Task Entry Point

What triggers this workflow? What is the desired output?

The workflow would be triggered by the closing of student feedback on a particular session. However, it wouldn't necessarily need to run immediately. For some sessions, we might want to see the feedback right away, but a brief overview might be enough. For end-of-course retrospective purposes, we might well want to run the LLM feedback module for each of the sets of session feedback and then maybe a more general version for all of them at once.

The eventual desired output would be a 1-2-page summary of the key themes and findings from student feedback on that session. What I have right now is an intermediate version that incorporates all student comments into themes so that it's easier to verify (or at the very least, spot check) the LLM's work.

Data Inputs

What information goes in? What is the data sensitivity classification?

Submitted data is a spreadsheet of student feedback on a particular course session. Identifying data for students (names, e-mail addresses) have been removed, and we'll also have to pseudonymize any mentions of the course instructors and other supporting faculty.

Since students submit this as a part of their course, I would say this counts as part of their "education record" and is hence FERPA-protected. As recommended by the [U of A AI Tools Guide](#), we'll be removing personally identifiable information as outlined above.

AI Tool Selected + Rationale

Why this tool for this task? Is it institutionally approved?

I decided at least initially to use Claude Sonnet 4.6 as accessed through the [U of A Open WebUI portal](#) as the initial tool for this task, though other LLMs would probably work. I chose this one because it's the highest-resource version we have access to, and I'm assuming that means it'll do the best job. However, I do wonder whether a lower resource version (or a different LLM in the portal) might do as good a job or even better. I think it would be worth exploring.

This tool is institutionally approved since it is provided through by the U of A.

Prompt Strategy

What prompting approach (zero-shot, few-shot, CoT)? Role specification?

I worked through refining the prompt with Claude before attempting the task. I am asking it to adopt a particular role for the task. Eventually, the goal will be for this to function as a zero-shot prompt, but additional refinement of the prompt will be necessary before we can get there. Perhaps a lot of refinement.

AI Output Review Checkpoint

What specifically do you check? For accuracy? Tone? Bias? Completeness?

I think there are two primary areas to check: (1) whether all student comments were used and were quoted accurately, and (2) whether the theming decisions the LLM made are reasonable.

1. There are enough student comments that checking them all each time this prompt is run would be unfeasible. I think a random sampling strategy might make sense: while we're testing it, take ~10% of the comments and see how they're reflected in the output. If there don't appear to be any issues, that would give me a bit more confidence that things are working as intended. This could also conceivably be done by another LLM prompt. If my report were to be generated in a file, asking an LLM to cross-reference students comments on the spreadsheet with what's in the generated report seems like it would be within its capabilities. I'm a little unsure whether using an LLM to check an LLM is clever or foolish, though. Perhaps a bit of both.
2. Checking the theming decisions would be a bit more difficult. One approach would be to go through a full qualitative theming process for one set of feedback and then compare those results to what the LLM generated. The problem is that there won't be one authoritative, correct version of how to do this - that's just how qualitative research works. If they're pretty close, then that would be a good sign. The second approach would be to read through the report to see if all the comments in the themes make sense. That would probably be the more sustainable method going forward. Though a third option would be to, as above, ask the LLM (or a different one, or in a different chat window) to check the theming and give feedback.

Human Verification Step

What human action confirms quality before the output is used?

See above. This is already covered in the AI Output Review Checkpoint.

Output + Disclosure Decision

Format, destination and disclosure action:

Ultimately, I would want to put these into a nicely formatted Microsoft Word document which could then be put into a folder shared between the course team. For disclosure to the course team, we'd just need to let them know. For disclosure to the students, it would need to either be put into the syllabus or mentioned in class.