

Output selected

I used an LLM-generated rejection letter from a journal that I used in [the discussion on Module 2 Task 2](#). I decided to use the shortest/smallest prompt because I expect that would carry with it more bias. (It's also possible that something like assigning it a role/persona might also carry bias, though.)

Representation

Whose perspectives are foregrounded? Whose are absent?

The perspectives of the journal and editorial team are foregrounded. The perspective of the submitting author is backgrounded. Perspectives from any other stakeholders are absent, or at best are indirectly referenced.

Population scaling

Which subgroups might receive systematically less useful outputs at scale?

This is difficult to answer without running the prompt many times and systematically analyzing the results. I don't think it's fair or reasonable to run a prompt once and rush to judgment. That said, I'd anticipate that the kinds of people who don't usually get to publish in this journal (or journals generally) are those who would be likely not to benefit. I'm not sure whether they would receive less useful outputs, though, since they're not writing the letters. And since the letter is being generated in support of a human-made decision (or the decision of a group of humans), I'm not sure there's a lot of risk for discrimination, etc. here.

Consequential decision

Which populations could be disadvantaged if this output informed institutional decisions?

I'm not sure how a form letter would have much of an impact on institutional decisions. If the decision on whether or not to accept or reject manuscripts was outsourced to an AI tool, that would potentially be a big problem. But if an LLM is just helping to speed along the process of drafting and sending out letters, I'm not sure there's a lot of harm done.

Analysis

150-word analysis and usage decision (use as-is / modify / use only with additional human review — and why):

From a bias and discrimination perspective, I think this output would be fine to use as-is. There might be some tweaking that needs to be done to bring it better in line with journal priorities, voice, etc., but the really important part here is that the editorial decision stays with the editorial board. The reading, processing, and decision-making processes (particularly the former two) do seem like they could be sped along with AI support, but that would run the risk of introducing rather more unseen bias into the mix. (Though, if I were to play devil's advocate for a moment, is that any more bias than a human would introduce?)