Section 1: Problem Statement

Session-level student feedback is underleveraged. Summaries and analyses of feedback for each session would make at-a-glance references easy, complementing faculty read-throughs. It could also be amalgamated into higher-level feedback to complement end-of-semester student feedback.

Section 2: Target Audience

There are two primary target audiences for this initiative: college instructional designers and college faculty. I would estimate AI fluency among this group to be at the moderately low to moderate level. I don't know that either group "needs" anything from this initiative, but both would stand to gain from having a quick, digestible reference to student feedback on each session. Major concerns are likely to be whether the LLM-generated reports are accurate and whether confidentiality of data is preserved.

Section 3: AI Tools and Approach

Identifying the specific tool to be used would likely be part of the testing process. However, the plan would be to only use models available through the U of A GenAI web portal (and thus pre-approved by the university). I initially developed the prompts for this task using Claude Opus 4.6, but I would be interested to see how other models would perform, particularly with an eye toward minimizing resource use by using lower-resource models. All models available through this platform are not trained off the data entered into them, and chats are only accessible to the user that created them. Since this is student feedback data gathered as a part of a course, it falls under FERPA.

Section 4: Ethical Safeguards

The key risks in this initiative all relate to confidentiality: student confidentiality and staff/faculty confidentiality. All of these will be the responsibility of whichever staff or faculty member is conducting this AI-enhanced analysis (hereafter "the analyst").

Student confidentiality: Before being uploaded for analysis, all student information (names and e-mail addresses) will be removed from the spreadsheet. This will be carried out by the analyst.

Staff/faculty confidentiality: Before being uploaded for analysis, all staff and faculty names will be changed to pseudonyms (e.g., “Dr. A,” “Dr. B,” etc.). This can be done efficiently through a find + replace operation in Microsoft Excel. The analyst who carries this out will need to ensure that all mentions of names are replaces, which may require a brief review of the comments. Once the analysis has been completed and moved into a document, the names can then be returned via another find + replace operation. Though not strictly necessary, this step is an extra precaution.

Section 5: Equity Considerations

I do not foresee any major equity concerns in this project.

Section 6: Success Metrics

The primary success metric would be whether a course team (or more than one) in the college takes up this process, trials it, and finds it useful such that they’d like to continue using it in the future.

In terms of accuracy and reliability, developing that process will be part of the rollout because I only had time to develop a prompt during this workshop. However, a successful prompt and tool would be one which:

Criterion	How to Assess
can reliably incorporate all student comments into themes	Verify that all student comments are included somewhere in the output. Could be done by hand by a human, but an LLM could likely accomplish this task very quickly (and likely well) as it’s a low-level task that requires processing a lot of data. Initial human verification would probably consist of randomly selecting 10% of student comments and verifying that they are in the output.
does not hallucinate any student comments which were not present in the original spreadsheet	Verify that no comments listed in the output lack a match in the uploaded spreadsheet. Could be done by hand by a human, but an LLM could likely accomplish this task very quickly (and likely well) as it’s a low-level task that requires processing a lot of data. Initial human verification would probably consist of randomly selecting 10% of student comments and verifying that they are in the output.

sorts students comments into reasonable, defensible themes (understanding that there is no one correct way to do so)	Two-phase human review: 1. Read over the raw spreadsheet data to get a feel for contours of the data. 2. Review the AI-generated output, paying attention to overall fit between comments and themes.
calculates numerical data averages correctly	Calculate averages by hand, which would require stripping the superfluous HTML formatting from the source spreadsheet.

One lingering question I have is whether these verifications need to be carried out every single time the process is carried out. In principle, the answer is "yes" because we don't want to just blindly trust LLM output. But practically, if we're going through this process dozens of times, and we do the checks ten times in a row and everything looks good, do we then relax the verification standards since we're somewhat confident we're getting an accurate and reliable product?

Section 7: Rollout Plan

Step	What	Who	Resource(s) Needed	By When
1	peer review and discussion of the prompt, task, output, and review process	me + peers (other instructional designers +/- Assessment, Evaluation & Analytics Team +/- interested faculty)	time	August 17
2	determine the "best" model for the job	me + peers (other instructional designers +/- Assessment, Evaluation & Analytics Team +/- interested faculty)	time	September 7
3	deliberate, intensive testing	me + peers (other instructional designers +/- Assessment, Evaluation & Analytics Team +/- interested faculty)	time	September 28
4	test of the prompt with real course material	me + interested course team members (faculty and/or instructional designers)	time	October 19

Elevator Pitch

We're working on using large language models to process anonymized student session feedback into a list of key themes as well as an executive summary. These will complement faculty review of this data and supplement end-of-semester feedback.